


Networks

Sridhar Alagar

How can two processes communicate?

- Both processes in the same machine
 - Use pipes, or other IPC mechanisms such a shared memory, message passing
- Both the processes are in different machines (that are far apart)
 - networks

Networking




Layers

- Computer has several layers
 - Physical, OS, libraries, application
- Similarly, network has layers
 - Each layer provide some abstraction


The Physical Layer

- Send/receive bits over a medium (wire/air)
 - 0 – off
 - 1 – on
- Bandwidth/Latency matters





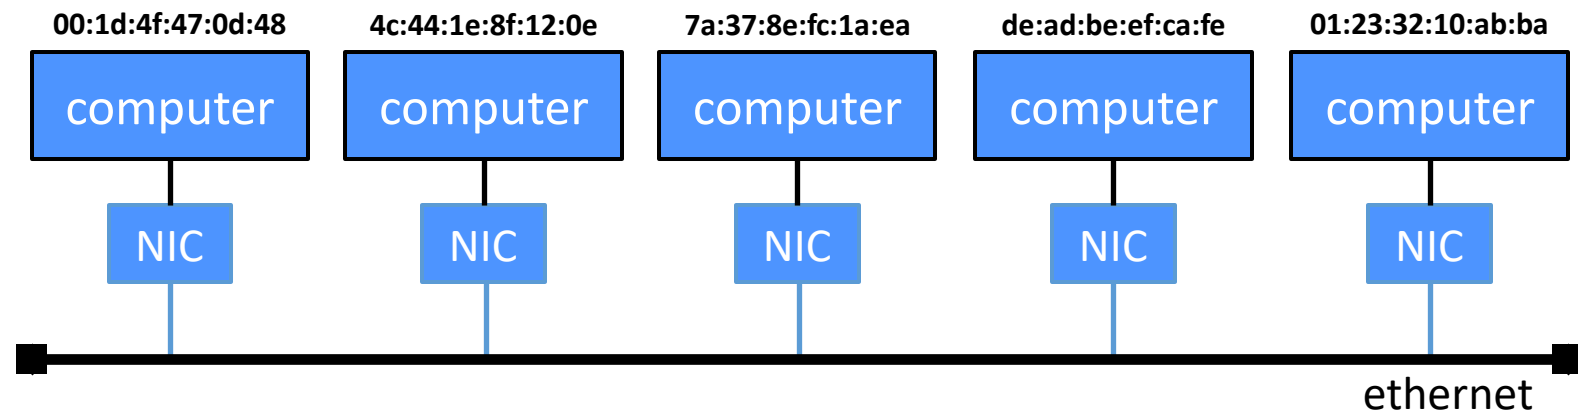
The Data Link Layer

- Link layer is responsible for sending packets between two computers that are directly connected
- Specifies how
 - bits are “packetized”
 - network interface controllers (NICs) are addressed




The Data Link Layer

- Multiple computers on a LAN contend for the network medium
 - Media access control (MAC) specifies how computers cooperate





The Network Layer

- Trying to transmit packets across different networks. This is where the “internet” is
- Individual networks are connected by routers that span networks





The Network Layer (IP)

- Internet Protocol (IP) routes packets across multiple networks
 - Individual networks are connected by routers that span networks
 - Every computer has a unique IP address



The Network Layer (IP)

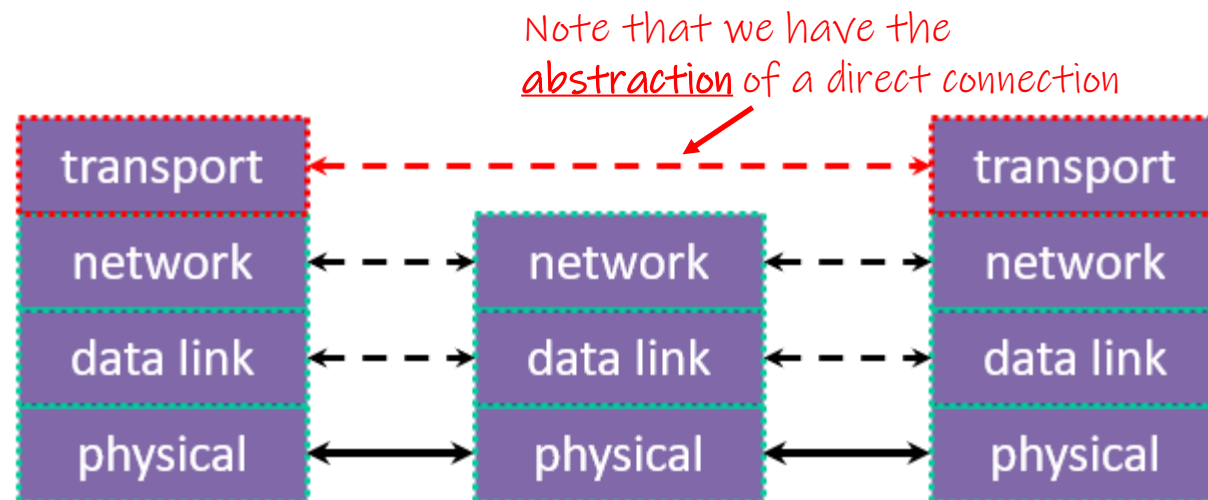
- Protocol to route packets
 - Best effort delivery



- Latency depends on
 - Distance
 - No of routers in path

The Transport Layer


- Provides an end-to-end (direct) connection
 - Hides the complexities of the network layer and below
- Provides different protocols to interface between source and destination
 - e.g., Transmission Control Protocol (TCP), User Datagram Protocol (UDP)
 - These protocols still work with packets, but manages their order, reliability, multiple applications using the network



The Transport Layer - TCP


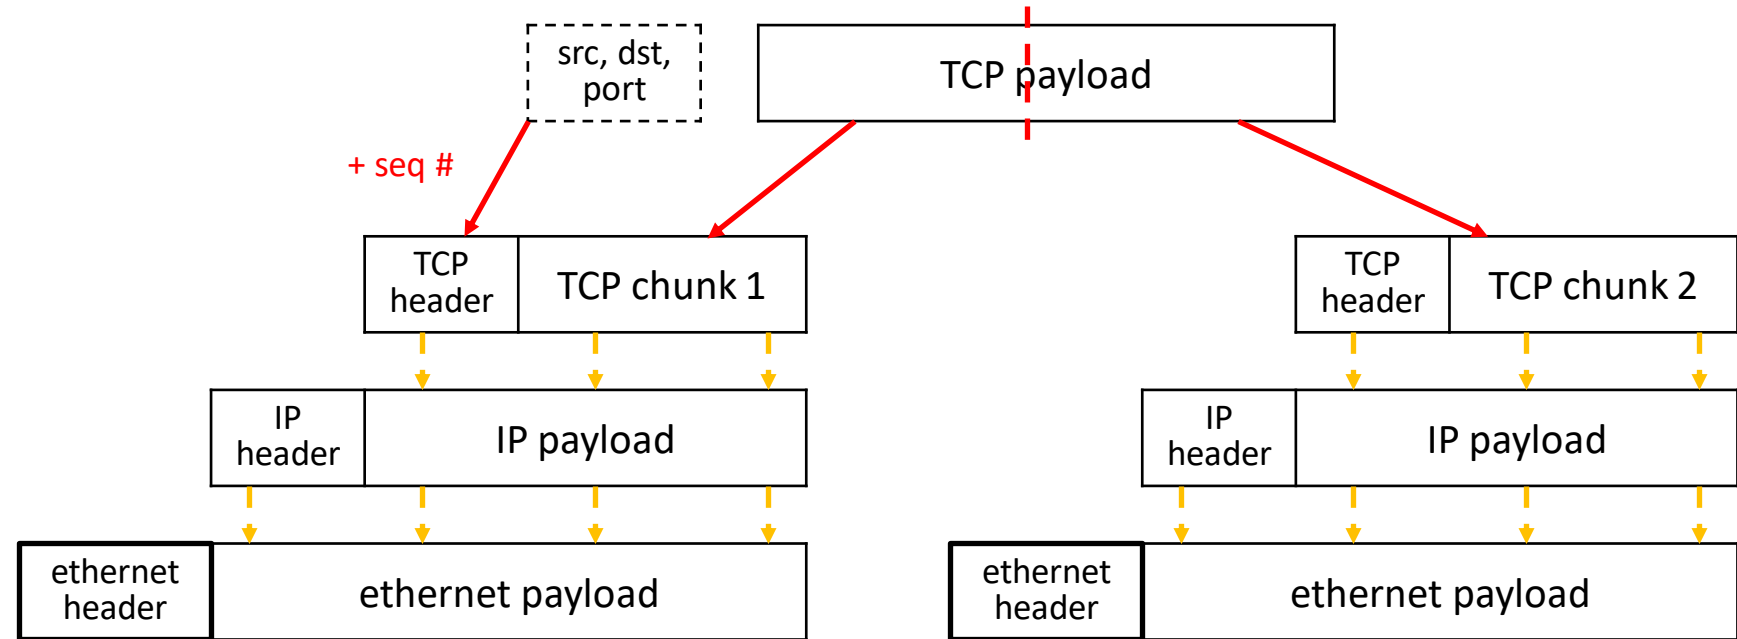
Transmission Control Protocol (TCP)

- Provides applications with reliable, ordered byte streams
 - Sends stream data as multiple IP packets (differentiated by sequence numbers) and retransmits them as necessary
 - When receiving, puts packets back in order and detects missing packets
- A single host (IP address) can have up to $2^{16} = 65,535$ “ports”
 - Kind of like an apartment number at a postal address (your applications are the residents who get mail sent to an apt. #)



The Transport Layer - TCP


- Packet encapsulation



The Transport Layer - UDP


User Datagram Protocol (UDP)

- Provides applications with unreliable packet delivery
- UDP is a thin, simple layer on top of IP
 - Datagrams still are fragmented into multiple IP packets




The (missing) Layers 5 and 6

- Layer 5 – session layer
 - Supposedly handles establishing and terminating application sessions
- Layer 6 – presentation layer
 - Datagrams still are fragmented into multiple IP packets



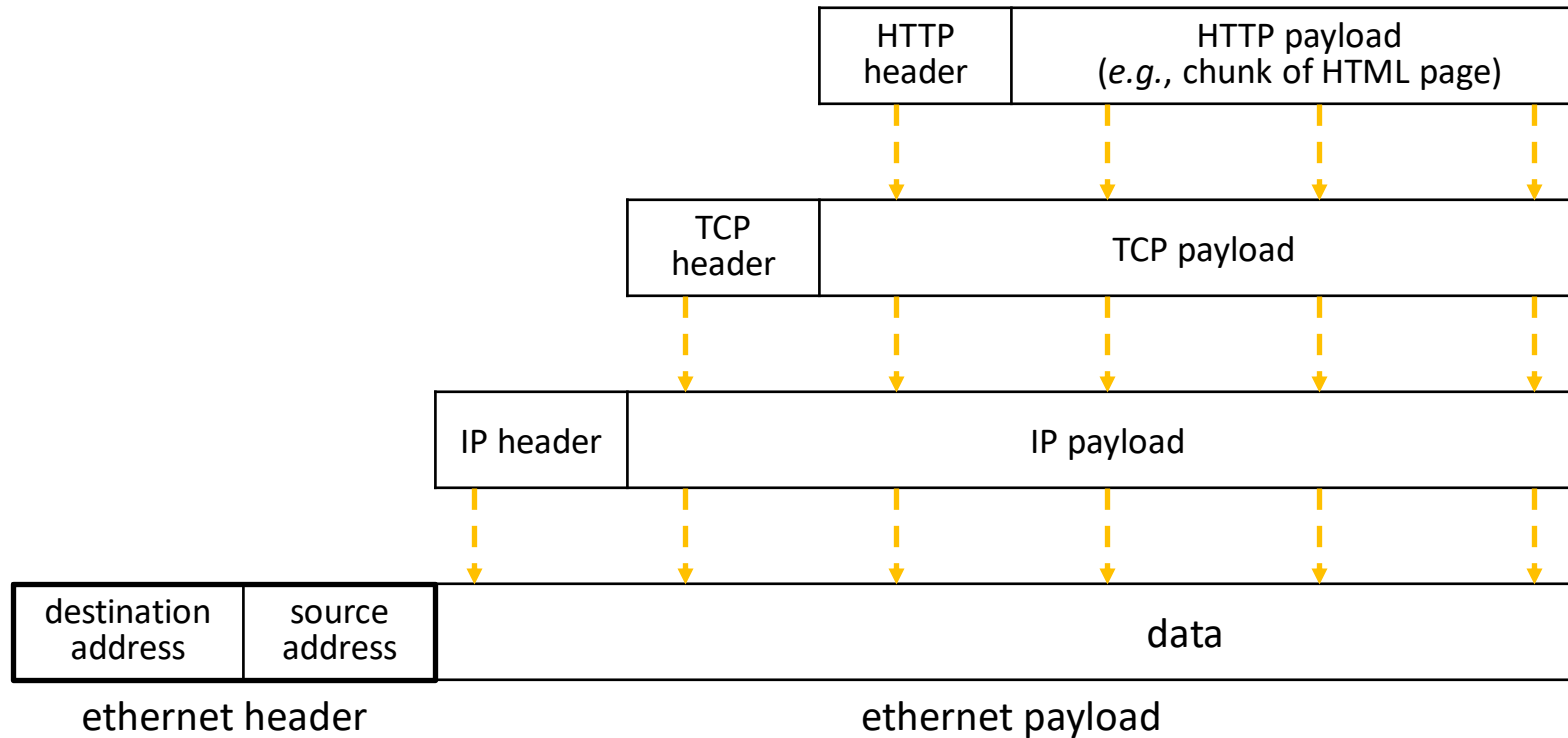
The Application Layer

- Application protocols
 - The format and meaning of messages between application entities
 - e.g., HTTP is an application-level protocol that dictates how web browsers and web servers communicate
 - HTTP is implemented on top of TCP streams




The Application Layer

- Packet encapsulation



The Application Layer

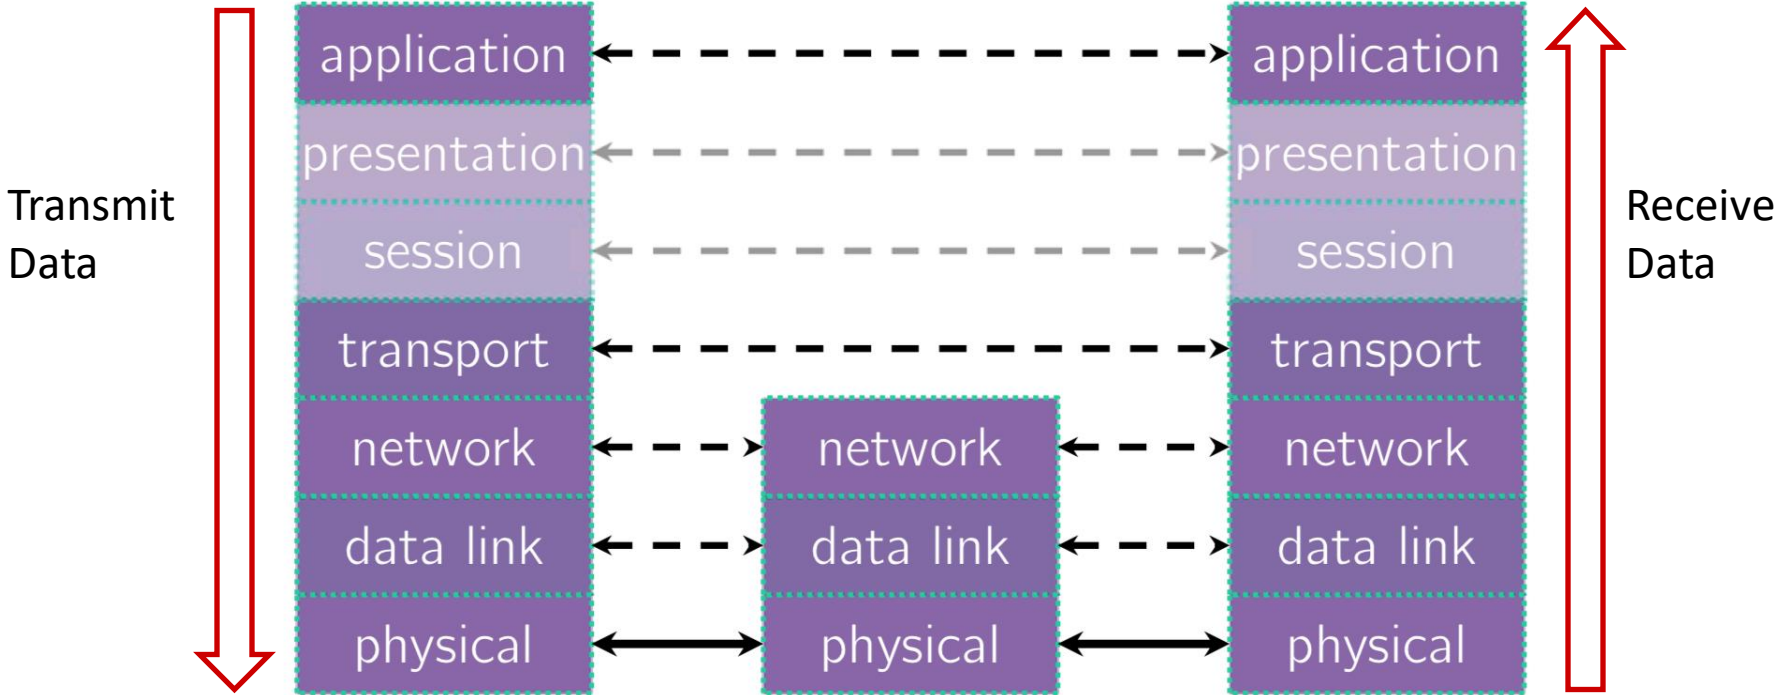
- Packet encapsulation



The Application Layer

- Popular application-level protocols:
- DNS: translates a domain name (e.g., `www.utdallas.edu`) into one or more IP addresses (e.g., `10.176.92.9`)
 - Domain Name System
 - A hierarchy of DNS servers cooperate to do this
- HTTP: web protocols
 - Hypertext Transfer Protocol
- SMTP, IMAP, POP: mail delivery and access protocols
 - Secure Mail Transfer Protocol, Internet Message Access Protocol, Post Office Protocol
- SSH: secure remote login protocol
 - Secure Shell
- bittorrent: peer-to-peer, swarming file sharing protocol

Actual Data Flow



Netcat demo

- netcat (`nc`) is “a computer networking utility for reading from and writing to network connections using TCP or UDP”
 - <https://en.wikipedia.org/wiki/Netcat>
 - Listen on port: `nc -l <port>`
 - Connect: `nc <hostname> <port>`


Files and File Descriptors

- Remember `open()`, `read()`, `write()`, and `close()`?
- POSIX system calls for interacting with files
 - `open()` returns a file descriptor
 - An integer that represents an open file
 - This file descriptor is then passed to `read()`, `write()`, and `close()`
- Inside the OS, the file descriptor is used to index into a table that keeps track of any OS-level state associated with the file, such as the file position (offset)
- Remember any other system call that returns file descriptor?

Network and Sockets

- UNIX likes to make *all* I/O look like file I/O
- You use **read** () and **write** () to communicate with remote computers over the network!
- A file descriptor used for network communications is called a **socket**
- Just like with files:
 - Your program can have multiple network channels open at once
 - You need to pass a file descriptor to **read** () and **write** () to let the OS know which network channel to use

File Descriptor Table



OS's File Descriptor Table for the Process

File Descriptor	Type	Connection
0	pipe	stdin (console)
1	pipe	stdout (console)
2	pipe	stderr (console)
3	TCP socket	local: 128.95.4.33:80 remote: 44.1.19.32:7113
5	file	index.html
8	file	pic.png
9	TCP socket	local: 128.95.4.33:80 remote: 102.12.3.4:5544

Types of Sockets

- Stream sockets
 - For connection-oriented, point-to-point, reliable byte streams
 - Using TCP, SCTP, or other stream transports
- Datagram sockets
 - For connection-less, one-to-many, unreliable packets
 - Using UDP or other packet transports
- Raw sockets
 - For layer-3 communication (raw IP packet manipulation)


Stream Sockets

- Typically used for client-server communications
 - **Client:** An application that establishes a connection to a server
 - **Server:** An application that receives connections from clients
 - Can also be used for other forms of communication like peer-to-peer

1) Establish connection:



2) Communicate:




3) Close connection:




Datagram Sockets

- Often used as a building block
 - No flow control, ordering, or reliability, so used less frequently
 - *e.g.* streaming media applications or DNS lookups

1) Create Sockets:



2) Communicate:



Sockets API

- It is the standard API for network programming
- Available on most OS
- Written in C

Socket API: Client TCP Connection

- We'll start by looking at the API from the point of view of a client connecting to a server over TCP
- There are five steps:
 - 1) Figure out the IP address and port to which to connect
 - 2) Create a socket
 - 3) Connect the socket to the remote server
 - 4) **read()** and **write()** data using the socket
 - 5) Close the socket

Step 1: Figure Out IP Address and Port

- Several parts:
 - Network addresses
 - Data structures for address info
 - DNS (Domain Name System) – finding IP addresses

IPv4 Addresses

- An IPv4 address is a **4-byte** tuple
 - For humans, written in “dotted-decimal notation”
 - *e.g.* 128.95.4.1 (80:5f:04:01 in hex)
- IPv4 address exhaustion
 - There are $2^{32} \approx 4.3$ billion IPv4 addresses
 - There are ≈ 7.87 billion people in the world (May 2021)

https://en.wikipedia.org/wiki/IPv4_address_exhaustion

IPv6 Addresses

- An IPv6 address is a **16-byte** tuple
 - Typically written in “hextets” (groups of 4 hex digits)
 - Can omit leading zeros in hextets
 - Double-colon replaces consecutive sections of zeros
 - *e.g.* `2d01:0db8:f188:0000:0000:0000:0000:1f33`
 - Shorthand: `2d01:db8:f188::1f33`
 - Transition is still ongoing


IPv4 Address Structures

```
#include <arpa/inet.h>
```

```
// IPv4 4-byte address
struct in_addr {
    uint32_t s_addr;           // Address in network byte order
};

// An IPv4-specific address structure
struct sockaddr_in {
    sa_family_t    sin_family; // Address family: AF_INET
    in_port_t      sin_port;   // Port in network byte order
    struct in_addr sin_addr;    // IPv4 address
    unsigned char  sin_zero[8]; // Pad out to 16 bytes
};
```

struct sockaddr_in:




IPv6 Address Structures

```
// IPv6 16-byte address
struct in6_addr {
    uint8_t s6_addr[16];           // Address in network byte order
};

// An IPv6-specific address structure
struct sockaddr_in6 {
    sa_family_t    sin6_family;    // Address family: AF_INET6
    in_port_t      sin6_port;      // Port number
    uint32_t        sin6_flowinfo; // IPv6 flow information
    struct in6_addr sin6_addr;      // IPv6 address
    uint32_t        sin6_scope_id; // Scope ID
};
```

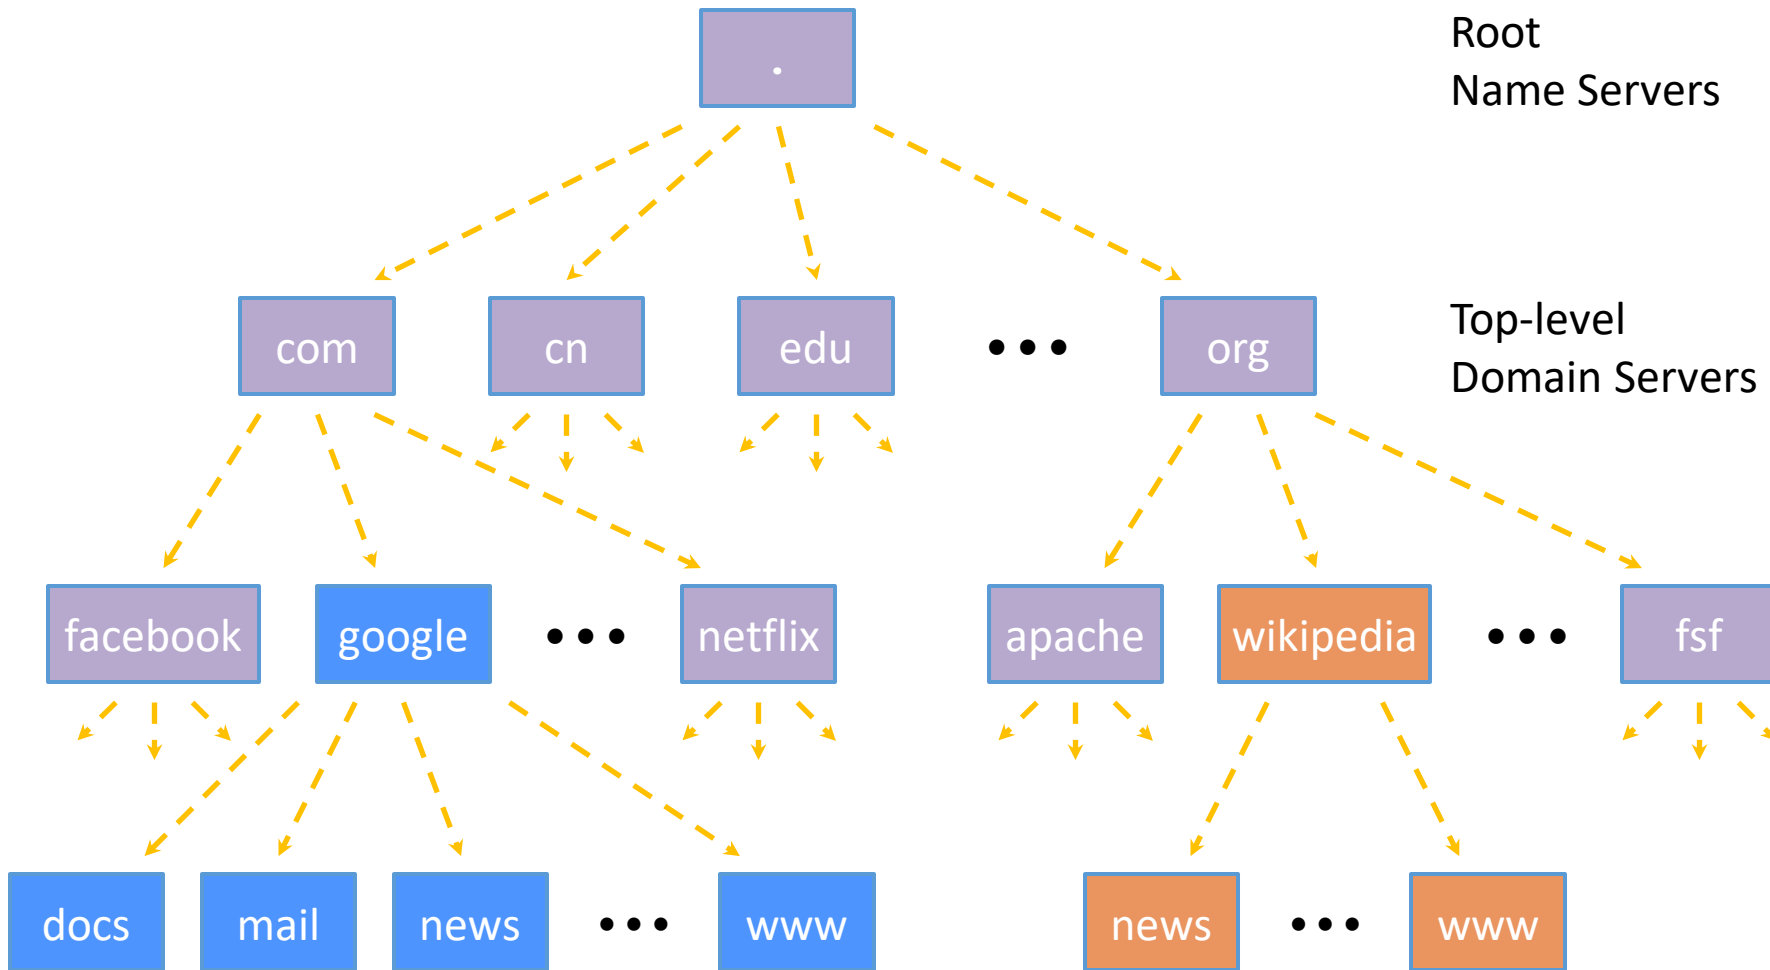
struct sockaddr_in6:



Domain Name System

- People tend to use DNS names, not IP addresses
 - The Sockets API lets you convert between the two
 - It's a complicated process, though:
 - A given DNS name can have many IP addresses
 - Many different IP addresses can map to the same DNS name
 - An IP address will reverse map into at most one DNS name
 - A DNS lookup may require interacting with many DNS servers
- You can use the Linux program “dig” to explore DNS
 - `dig @server name type (+short)`
 - `server`: specific name server to query
 - `type`: A (IPv4), AAAA (IPv6), ANY (includes all types)

DNS Hierarchy



Resolving DNS Names

- Use sys call **getaddrinfo()** to get ip addresses for a name

```
int getaddrinfo(const char* hostname,
                const char* service,
                const struct addrinfo* hints,
                struct addrinfo** res);
```

- Basic idea:
 - Tell **getaddrinfo()** which host and port you want resolved
 - String representation for host: DNS name or IP address
 - Set up a “hints” structure with constraints you want respected
 - **getaddrinfo()** gives you a list of addresses packed into an “addrinfo” structure/linked list
 - Returns **0** on success; returns *negative number* on failure
 - Free the `struct addrinfo` later using **freeaddrinfo()**

DNS Lookup Procedure

```
struct addrinfo {  
    int      ai_flags;           // additional flags  
    int      ai_family;         // AF_INET, AF_INET6, AF_UNSPEC  
    int      ai_socktype;       // SOCK_STREAM, SOCK_DGRAM, 0  
    int      ai_protocol;       // IPPROTO_TCP, IPPROTO_UDP, 0  
    size_t   ai_addrlen;        // length of socket addr in bytes  
    struct sockaddr* ai_addr;    // pointer to socket addr  
    char*     ai_canonname;      // canonical name  
    struct addrinfo* ai_next;    // can form a linked list  
};
```

- 1) Create a `struct addrinfo hints`
 - 2) Zero out hints for “defaults”
 - 3) Set specific fields of hints as desired
 - 4) Call `getaddrinfo()` using `&hints`
 - 5) Resulting linked list `res` will have all fields appropriately set
- See `dnsresolve.c`

Socket API: Client TCP connection

- There are five steps:
 - 1) Figure out the IP address and port to connect to
 - 2) Create a socket
 - 3) Connect the socket to the remote server
 - `read()` and `write()` data using the socket
 - 5) Close the socket

Step 2: Creating a Socket

- Use the **socket** () system call

```
int socket(int domain, int type, int protocol);
```

- Creates an endpoint for communication
- Creating a socket doesn't bind it to a local address or port yet
- Returns file descriptor or -1 on error

Step 3: Connect to a Server

- The **connect** () system call establishes a connection to a remote host

```
int connect(int sockfd, const struct sockaddr* addr,  
            socklen_t addrlen);
```

- sockfd: Socket file description from Step 2
 - addr and addrlen: Usually from one of the address structures returned by getaddrinfo in Step 1 (DNS lookup)
 - Returns 0 on success and -1 on error
- **connect** () may take some time to return
 - It is a *blocking* call by default
 - The network stack within the OS will communicate with the remote host to establish a TCP connection to it
 - This involves ~2 *round trips* across the network

Step 4: read()

- If there is data that has already been received by the network stack, then read will return immediately with it
 - `read()` might return with *less* data than you asked for
- If there is no data waiting for you, by default `read()` will *block* until something arrives
 - Can `read()` return 0?

Step 4: write()

- **write** () queues your data in a send buffer in the OS and then returns
 - The OS transmits the data over the network in the background
 - When **write** () returns, the receiver probably has not yet received the data!
- If there is no more space left in the send buffer, by default **write** () will *block*

Step 5: close()

- Straightforward. Same function as with file I/O.
- Shuts down the socket and frees resources and file descriptors associated with it on both ends of the connection

Summary: Client TCP Connection

- There are five steps:
 - 1) Figure out the IP address and port to which to connect
 - 2) Create a socket
 - 3) Connect the socket to the remote server
 - 4) **read()** and **write()** data using the socket
 - 5) Close the socket

Socket API: Server TCP Connection

- Pretty similar to clients, but with additional steps:
 - 1) Figure out the IP address and port on which to listen
 - 2) Create a socket
 - 3) **bind()** the socket to the address(es) and port
 - 4) Tell the socket to **listen()** for incoming clients
 - 5) **accept()** a client connection
read() and **write()** to that connection
 - 7) **close()** the client socket

Step 1: Figure out IP address(es) & Port

- **getaddrinfo** () invocation may or may not be needed (but we'll use it)
 - Do you know your IP address(es) already?
 - Static vs. dynamic IP address allocation
 - Even if the machine has a static IP address, don't wire it into the code – either look it up dynamically or use a configuration file
 - Can request listen on all local IP addresses by passing `NULL` as `hostname` and setting `AI_PASSIVE` in `hints.ai_flags`

Step 2: Create a Socket

- **socket** () call is same as before

```
int socket(int domain, int type, int protocol);
```

- Can directly use constants or fields from result of **getaddrinfo** ()
- Recall that this just returns a file descriptor – IP address and port are not associated with socket yet

Step 3: Bind the socket

- ```
int bind(int sockfd, const struct sockaddr* addr,
 socklen_t addrlen);
```

- Looks nearly identical to **connect()** !
- Returns 0 on success, -1 on error

# Step 4: Listen for Incoming Clients

- ```
int listen(int sockfd, int backlog);
```

 - Tells the OS that the socket is a listening socket that clients can connect to
 - `backlog`: maximum length of connection queue
 - Gets truncated, if necessary, to defined constant `SOMAXCONN`
 - The OS will refuse new connections once queue is full until server `accept()` s them (removing them from the queue)
 - Returns `0` on success, `-1` on error
 - Clients can start connecting to the socket as soon as `listen()` returns
 - Server can't use a connection until you `accept()` it

Step 5: Accept a Client Connection

- ```
int accept(int sockfd, struct sockaddr* addr,
 socklen_t* addrlen);
```
- Returns an active, ready-to-use socket file descriptor connected to a client (or `-1` on error)
  - `sockfd` must have been created, bound, *and* listening
  - Pulls a queued connection or waits for an incoming one
- `addr` and `addrlen` are output parameters
  - `*addrlen` should initially be set to `sizeof(*addr)`, gets overwritten with the size of the client address
  - Address information of client is written into `*addr`
    - Use `inet_ntop()` to get the client's printable IP address
    - Use `getnameinfo()` to do a *reverse DNS lookup* on the client

# Example

- See `server.c`
  - Takes in a port number from the command line
  - Opens a server socket, prints info, then listens for connections
    - *Can connect to it using netcat (`nc`)*
  - Accepts connections as they come
  - Echoes any data the client sends to it on `stdout` and also sends it back to the client